

---

## Lecture 8 In-Class Example *Solutions*

---

**Example 1.** This example is taken from *Mastering Metrics* chapter 4 and is based on Carpenter & Dobkin (2009). The analysis revisits the question of whether legal access to alcohol is associated with higher mortality rates. The dataset referenced below includes death rates by age, in which age (from 19-23) is divided into 50 equal-width cells. Individuals are considered “treated” when they reach 21 and can legally drink alcohol. The discontinuity in treatment is sharp, since *all* persons who reach this age are treated. If legal access to alcohol is associated with higher mortality, we would expect to see a discontinuity in death rates at 21. Read the file *AEJfigs.dta* into Stata using:


```
use https://github.com/spcorcor18/LP0-8852/raw/main/data/AEJfigs.dta
```

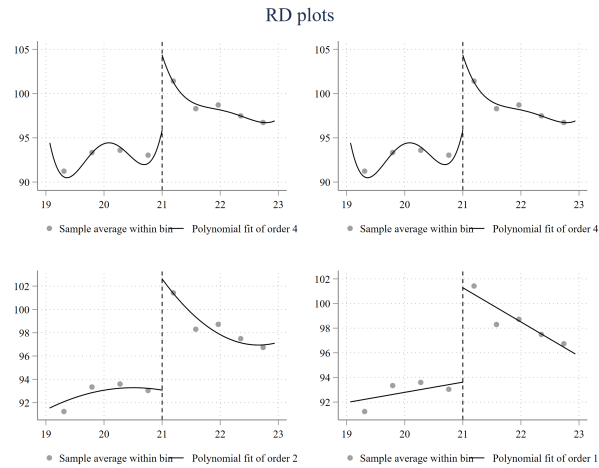
- (a) Create a new age variable *age* that is centered at 21, and a “treatment” variable *over21* that equals 1 if *age* > 0 (and 0 otherwise).

```
gen age = agecell - 21  
gen over21 = agecell >= 21
```

- (b) Create a scatterplot of mortality rates from all cases (*all*) against *age*. Is a discontinuity evident? Try a few RD plots using different bin methods (e.g., ES or QS, IMSE or MV) and different polynomial orders (e.g., quartic (the default), quadratic, linear). Note the RD plot does not have a lot of added value here, since the data are already binned.

**The scatterplot and RD plots are shown below. In all five plots a discontinuity is evident at age 21. Again, since the data are already binned, the scatterplot by itself is quite revealing of the discontinuity.**



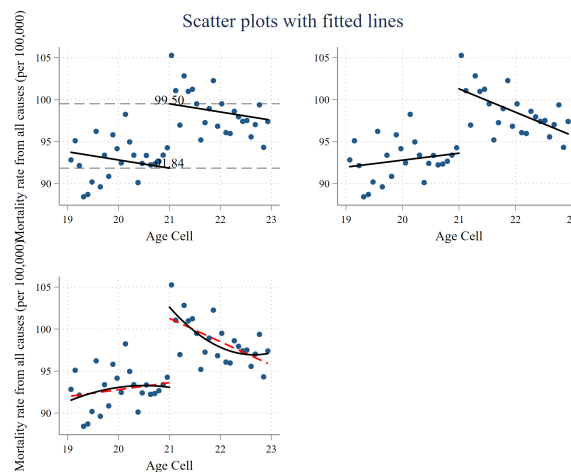


(c) We will eventually use `rdrobust` for estimation. But for now use OLS regression to estimate the RD under different scenarios:

- Linear model, assuming the same slope on both sides
- Linear model, assuming different slopes on each side
- Quadratic model, assuming the same slope terms on both sides
- Quadratic model, assuming different slope terms on each side

How are the coefficients interpreted in these models?

Get predicted values and overlay these predicted values on the scatterplot. (Note: this is what `rdplot` can do for you, but this is for illustration). You can easily reproduce Figure 4.2 in *Mastering Metrics* (linear model, same slopes).



Regression results shown below. In the first model, the slope on age is assumed to be the same on both sides of 21. This is estimated to be -0.975, meaning all-cause mortality is declining with age. The intercept shift at 21 is 7.66, meaning there is a jump in all-causes mortality at age 21 of 7.66. Note that *age* is centered at 21, so the constant term of 91.8 is the predicted all-cause mortality rate when *age*=0 (i.e., age 21)

In the second model, the close on age is allowed to differ on each side of 21. The slope below 21 is estimated to be 0.827; above 21 the slope is -3.60 *lower* (the interaction term). The discontinuity at age 21 is a very similar 7.66.

The third model is a quadratic with the linear and quadratic terms slopes assumed to be the same on each side of 21. In the fourth model, the slopes are allowed to differ on each side of 21. The interaction terms indicate how *different* the slope terms are above 21 vs. below. The discontinuity in the third model is 7.66; in the fourth it is a larger 9.55.

```
.          // linear model in age with intercept shift at over21
.          reg all age over21
```

Source	SS	df	MS	Number of obs	=	
Model	410.138151	2	205.069075	F(2, 45)	=	32.99
Residual	279.682408	45	6.21516463	Prob > F	=	0.0000
				R-squared	=	0.5946
				Adj R-squared	=	0.5765
Total	689.820559	47	14.6770332	Root MSE	=	2.493

	all	Coef.	Std. Err.	t	P> t	[95% Conf. Interval]
age		-.9746843	.6324613	-1.54	0.130	-2.248527 .2991581
over21		7.662709	1.440286	5.32	0.000	4.761824 10.56359
_cons		91.84137	.8050394	114.08	0.000	90.21994 93.4628

```
.          predict allfitlin
(option xb assumed; fitted values)

.          // linear model in age with intercept shift at over21, different slope
.          // below and above c
.          reg all c.age##i.over21
```

Source	SS	df	MS	Number of obs	=	
Model	460.574058	3	153.524686	F(3, 44)	=	29.47
Residual	229.246501	44	5.21014775	Prob > F	=	0.0000
				R-squared	=	0.6677
				Adj R-squared	=	0.6450

Total | 689.820559      47   14.6770332    Root MSE      =      2.2826

all	Coef.	Std. Err.	t	P> t	[95% Conf. Interval]	
age	.8269952	.8189316	1.01	0.318	-.823453	2.477443
1.over21	7.662709	1.318704	5.81	0.000	5.005035	10.32038
over21#c.age						
1	-3.603359	1.158144	-3.11	0.003	-5.937445	-1.269273
_cons	93.61837	.9324647	100.40	0.000	91.73911	95.49763

```
. predict allfitlini
(option xb assumed; fitted values)
```

```
. reg all c.age##c.age over21
```

Source	SS	df	MS	Number of obs	=	48
Model	453.339903	3	151.113301	F(3, 44)	=	28.12
Residual	236.480656	44	5.37456036	Prob > F	=	0.0000
Total	689.820559	47	14.6770332	R-squared	=	0.6572
				Adj R-squared	=	0.6338
				Root MSE	=	2.3183

all	Coef.	Std. Err.	t	P> t	[95% Conf. Interval]	
age	-.9746843	.5881378	-1.66	0.105	-2.159998	.2106296
c.age#c.age	-.8186505	.2887482	-2.84	0.007	-1.400584	-.2367167
over21	7.662709	1.339349	5.72	0.000	4.963428	10.36199
_cons	92.90274	.8370061	110.99	0.000	91.21587	94.58962

```
. predict allfitq
(option xb assumed; fitted values)
```

```
. // quadratic model in age with intercept shift at over21, different slopes
. // below and above c
. reg all c.age##c.age##i.over21
```


Source	SS	df	MS	Number of obs	=	48
Model	470.512103	5	94.1024205	F(5, 42)	=	18.02
Residual	219.308457	42	5.22162992	Prob > F	=	0.0000
Total	689.820559	47	14.6770332	R-squared	=	0.6821
				Adj R-squared	=	0.6442
				Root MSE	=	2.2851

	all	Coef.	Std. Err.	t	P> t	[95% Conf. Interval]
	age	-.8305827	3.290064	-0.25	0.802	-7.470202 5.809036
	c.age#c.age	-.8402999	1.615268	-0.52	0.606	-4.100043 2.419443
	1.over21	9.547789	1.985277	4.81	0.000	5.541337 13.55424
	over21#c.age					
	1	-6.017014	4.652854	-1.29	0.203	-15.40685 3.372825
	over21#c.age#					
	c.age					
	1	2.904189	2.284334	1.27	0.211	-1.705784 7.514162
	_cons	93.07294	1.403803	66.30	0.000	90.23995 95.90593

```
. predict allfitqi
```

- (d) Repeat the linear and quadratic models for two other measures of mortality: deaths by motor vehicle accidents and deaths by internal causes. One might expect the former to be most affected by access to alcohol. The latter could be considered a placebo test. You can easily reproduce Figure 4.5 in *Mastering Metrics* that combines these.

**Figure is shown below. The discontinuity is evident for motor vehicle accidents but not for deaths by internal causes.**



- (e) Try part (c) using `rdrobust`. For comparability, use a uniform kernel and a bandwidth of 2. Compare the point estimates and standard errors to (c).

**Results shown below. The point estimates of 7.66 and 9.55 correspond to those in part (c) in which we allowed the linear and quadratic terms to vary**

on either side of age 21. We did not ask for the optimal bandwidth here (the manual bandwidth  $h=2$ ), so `rdrobust` used all of the data, 24 points on each side. For comparability with (c), we did not use weighting here, but rather a uniform kernel.

```
.      rdrobust all age, c(0) h(2) kernel(uniform) p(1)
```

Sharp RD estimates using local polynomial regression.

Cutoff c = 0   Left of c   Right of c			Number of obs = 48	
-----+-----			BW type = Manual	
Number of obs	24	24	Kernel =	Uniform
Eff. Number of obs	24	24	VCE method =	NN
Order est. (p)	1	1		
Order bias (q)	2	2		
BW est. (h)	2.000	2.000		
BW bias (b)	2.000	2.000		
rho (h/b)	1.000	1.000		

Outcome: all. Running variable: age.

Method	Coef.	Std. Err.	z	P> z	[95% Conf. Interval]	
-----+-----						
Conventional	7.6627	1.3089	5.8541	0.000	5.09721	10.2282
Robust	-	-	4.3630	0.000	5.25872	13.8369
-----+-----						

```
.      rdrobust all age, c(0) h(2) kernel(uniform) p(2)
```

Sharp RD estimates using local polynomial regression.

Cutoff c = 0   Left of c   Right of c			Number of obs =		48
-----+-----			BW type =		Manual
Number of obs	24	24	Kernel =		Uniform
Eff. Number of obs	24	24	VCE method =		NN
Order est. (p)	2	2			
Order bias (q)	3	3			
BW est. (h)	2.000	2.000			
BW bias (b)	2.000	2.000			
rho (h/b)	1.000	1.000			

Outcome: all. Running variable: age.

Method	Coef.	Std. Err.	z	P> z	[95% Conf. Interval]	
-----+-----						
Conventional	9.5478	2.1883	4.3630	0.000	5.25872	13.8369
Robust	-	-	3.6386	0.000	5.19606	17.3296
-----+-----						

(f) Use `rdbwselect` to find the optimal MSE bandwidth under different choices of polyno-

mial (linear or quadratic) and kernel (uniform or triangular). Use these in `rdrobust` and compare your point estimates and standard errors. Normally, one could pass through the optimal bandwidth to `rdplot`, but this is problematic with these data (they are already binned).

Four optimal bandwidths are found below. In each case we asked for the bandwidth to be the same width on the left and right of age 21 (`mserd`). The four bandwidths are 0.493, 0.698, 0.451, and 0.752. These correspond to the linear and quadratic models (triangular kernel), and linear and quadratic models (uniform kernel). In the last step we put the `bwselect` right in the `rdrobust` command so that Stata uses the optimal bandwidth, in this case for a quadratic model with triangular kernel. The result is 8 observations per side and a discontinuity estimate of 10.28. Again, these data are already binned, which is why the resulting number of observations is so small.

```
.      rdbwselect all age, c(0) bwselect(mserd) kernel(triangular) p(1)
```

Bandwidth estimators for sharp RD local polynomial regression.

Cutoff c =	Left of c	Right of c	Number of obs =	48
-----+-----			Kernel	= Triangular
Number of obs	24	24	VCE method	= NN
Min of age	-1.932	0.041		
Max of age	-0.041	1.932		
Order est. (p)	1	1		
Order bias (q)	2	2		

Outcome: all. Running variable: age.

	BW est. (h)		BW bias (b)	
Method	Left of c	Right of c	Left of c	Right of c
-----+-----+-----				
mserd	0.493	0.493	0.780	0.780
-----+-----+-----				

```
.      rdbwselect all age, c(0) bwselect(mserd) kernel(triangular) p(2)
```

Bandwidth estimators for sharp RD local polynomial regression.

Cutoff c =	Left of c	Right of c	Number of obs =	48
-----+-----			Kernel	= Triangular
Number of obs	24	24	VCE method	= NN
Min of age	-1.932	0.041		
Max of age	-0.041	1.932		
Order est. (p)	2	2		
Order bias (q)	3	3		

Outcome: all. Running variable: age.

Method	BW est. (h)		Method	BW bias (b)	
	Left of c	Right of c		Left of c	Right of c
mserd	0.698	0.698		0.961	0.961

```
. rdbwselect all age, c(0) bwselect(mserd) kernel(uniform) p(1)
```

Bandwidth estimators for sharp RD local polynomial regression.

Cutoff c =	Left of c	Right of c	Number of obs =	48
			Kernel =	Uniform
			VCE method =	NN
Number of obs	24	24		
Min of age	-1.932	0.041		
Max of age	-0.041	1.932		
Order est. (p)	1	1		
Order bias (q)	2	2		

Outcome: all. Running variable: age.

Method	BW est. (h)		Method	BW bias (b)	
	Left of c	Right of c		Left of c	Right of c
mserd	0.451	0.451		0.711	0.711

```
. rdbwselect all age, c(0) bwselect(mserd) kernel(uniform) p(2)
```

Bandwidth estimators for sharp RD local polynomial regression.

Cutoff c =	Left of c	Right of c	Number of obs =	48
			Kernel =	Uniform
			VCE method =	NN
Number of obs	24	24		
Min of age	-1.932	0.041		
Max of age	-0.041	1.932		
Order est. (p)	2	2		
Order bias (q)	3	3		

Outcome: all. Running variable: age.

Method	BW est. (h)		Method	BW bias (b)	
	Left of c	Right of c		Left of c	Right of c
mserd	0.752	0.752		1.062	1.062

```
. // Pass through the optimal bandwidth to rdrobust and rdplot
. // (just using the quadratic, triangular case)
```

```
. rdrobust all age, c(0) bwselect(mserd) kernel(triangular) p(2)
```

Estimates might be unreliable due to low number of effective observations.



Sharp RD estimates using local polynomial regression.

Cutoff c = 0   Left of c    Right of c			Number of obs =	48
-----+-----			BW type =	mserd
Number of obs	24	24	Kernel =	Triangular
Eff. Number of obs	8	8	VCE method =	NN
Order est. (p)	2	2		
Order bias (q)	3	3		
BW est. (h)	0.698	0.698		
BW bias (b)	0.961	0.961		
rho (h/b)	0.726	0.726		

Outcome: all. Running variable: age.

-----							
Method	Coef.	Std. Err.	z	P> z	[95% Conf. Interval]		
-----+-----							
Conventional	10.28	4.7765	2.1522	0.031	.918045	19.6417	
Robust	-	-	1.9408	0.052	-.106652	21.7024	
-----							

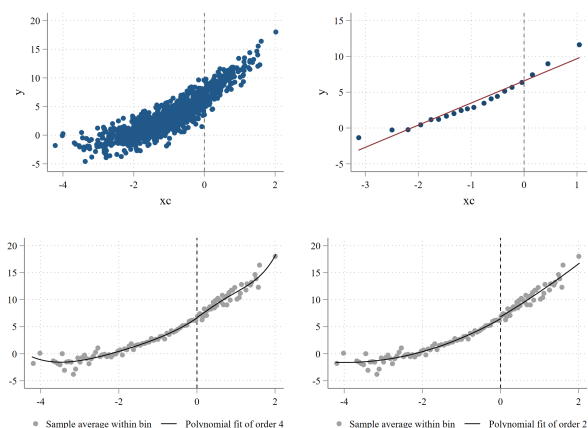
**Example 2.** This example will generate data with a known discontinuity in  $y$  at a threshold level of  $x$ , and then estimate a RD model. It illustrates various RD commands and results, and simulates manipulation of the running variable. (Adapted from Dale Ballou).

- (a) First produce simulated data using the syntax below. Notice that  $x$  is the running variable. What is the functional relationship between the outcome  $y$  and the running variable? What is the cut score? What is the treatment effect? Is this a sharp or fuzzy regression discontinuity? Create a new variable  $xc$  that is  $x$  centered at the cut score.

```
clear
set seed 1234
drawnorm x w e u, n(1000)
gen y = 3 + 3*x + .5*x^2 + w + u
gen t = (x > 1)
replace y = y + .5*t
// centered version of x--use this below
gen xc = x-1
```

There is a quadratic relationship between the running variable  $x$  and the outcome  $y$ , as seen in line 4. The cut score is  $x=1$  (line 5). The treatment effect is 0.5 (line 6): cases where  $t=1$  have a value of  $y$  that is 0.5 higher than what it would be otherwise. This is a *strict* regression discontinuity since all cases where  $x \leq 1$  are untreated, and all cases where  $x > 1$  are treated.

- (b) Produce a scatterplot of  $y$  against  $xc$ . Do you see evidence of a discontinuity? Try using `binscatter` and `rdplot`, the latter implementing a quartic  $p = 4$  and then quadratic  $p = 2$ . Do you see a discontinuity in these plots? Which shows it best?



Figures shown above. It is difficult to see any discontinuity in the scatter plot. The discontinuity in the binned scatter and RD plots is evident, but slight.

- (c) Estimate several (global, parameter) RD models using OLS, below. Do not use a kernel and use the full range of data. How close do these get to estimating the true treatment effect? Which model performs best, and why?

- Linear model, assuming the same slope on both sides
- Linear model, assuming different slopes on each side
- Quadratic model, assuming the same slope terms on both sides
- Quadratic model, assuming different slope terms on each side

Results below. The treatment effects in the two linear models are 2.1 and 1.1, quite a bit larger than the known effect of 0.5. The reason is that the functional relationship between  $y$  and  $x$  is misspecified. It is known to be quadratic, and we fit a linear model. The increasing slope of the relationship between  $y$  and  $x$  is mistakenly subsumed into the treatment effect. The treatment effects in the quadratic models are 0.45 and 0.51, much closer to the known effect of 0.5.

```
. reg y xc t
```

Source	SS	df	MS	Number of obs	=	1,000
--------	----	----	----	---------------	---	-------

-----+-----					F(2, 997)	=	2240.02
Model		10240.5295	2	5120.26474	Prob > F	=	0.0000
Residual		2278.9563	997	2.28581374	R-squared	=	0.8180
-----+-----					Adj R-squared	=	0.8176
Total		12519.4858	999	12.5320178	Root MSE	=	1.5119

-----+-----						
y		Coef.	Std. Err.	t	P> t	[95% Conf. Interval]
-----+-----						
xc		2.598992	.0622848	41.73	0.000	2.476767 2.721216
t		2.070868	.1711322	12.10	0.000	1.735047 2.406689
_cons		5.752373	.0962028	59.79	0.000	5.56359 5.941157
-----+-----						

. reg y xc t c.xc#c.xc

Source		SS	df	MS	Number of obs	=	1,000
-----+-----					F(3, 996)	=	1761.62
Model		10534.1874	3	3511.3958	Prob > F	=	0.0000
Residual		1985.29838	996	1.99327146	R-squared	=	0.8414
-----+-----					Adj R-squared	=	0.8409
Total		12519.4858	999	12.5320178	Root MSE	=	1.4118

-----+-----						
y		Coef.	Std. Err.	t	P> t	[95% Conf. Interval]
-----+-----						
xc		4.033263	.1317049	30.62	0.000	3.774812 4.291714
t		.4473678	.2083961	2.15	0.032	.0384221 .8563135
c.xc#c.xc		.5029089	.0414335	12.14	0.000	.4216019 .5842158
_cons		6.429792	.105761	60.80	0.000	6.222252 6.637332
-----+-----						

. reg y c.xc##i.t

Source		SS	df	MS	Number of obs	=	1,000
-----+-----					F(3, 996)	=	1612.85
Model		10382.3149	3	3460.77165	Prob > F	=	0.0000
Residual		2137.17084	996	2.14575385	R-squared	=	0.8293
-----+-----					Adj R-squared	=	0.8288
Total		12519.4858	999	12.5320178	Root MSE	=	1.4648

-----+-----						
y		Coef.	Std. Err.	t	P> t	[95% Conf. Interval]
-----+-----						
xc		2.486569	.0619109	40.16	0.000	2.365079 2.60806
1.t		1.132005	.2020684	5.60	0.000	.7354766 1.528534
t#c.xc						
1		2.25286	.2771458	8.13	0.000	1.709003 2.796716

_cons		5.60662	.0949179	59.07	0.000	5.420358	5.792882
-----							
. reg y c.xc##c.xc##i.t							
Source		SS	df	MS	Number of obs	=	1,000
-----					F(5, 994)	=	1055.08
Model		10534.5475	5	2106.90949	Prob > F	=	0.0000
Residual		1984.93832	994	1.99691984	R-squared	=	0.8415
-----					Adj R-squared	=	0.8407
Total		12519.4858	999	12.5320178	Root MSE	=	1.4131
-----							
y		Coef.	Std. Err.	t	P> t	[95% Conf. Interval]	
-----							
xc		4.055002	.18967	21.38	0.000	3.682803	4.427202
c.xc#c.xc		.5096664	.0584983	8.71	0.000	.3948719	.6244608
1.t		.3831426	.2649436	1.45	0.148	-.1367704	.9030556
t#c.xc							
1		.2596287	.8106604	0.32	0.749	-1.331174	1.850431
t#c.xc#c.xc							
1		-.2146489	.5198822	-0.41	0.680	-1.234842	.8055438
_cons		6.442093	.13259	48.59	0.000	6.181905	6.702282
-----							

- (d) Obtain local, nonparametric RD estimates using `rdrobust`. Use a linear fit, the optimal MSE bandwidth selection, and triangular kernel. Pass through the optimal bandwidth to `rdplot` to get a local RD plot. Note: the accompanying do-file also shows how to use the older command `rd`. See the help menu for that command for options.

**Results shown below. The point estimate using the optimal bandwidth, triangular kernel, and linear model is 0.41 with a confidence interval of (-0.52, 1.35).**

```
. rdrobust y xc, c(0) p(1) bwselect(mserd) kernel(triangular)
```

Sharp RD estimates using local polynomial regression.

Cutoff c = 0		Left of c	Right of c	Number of obs	=	1000
-----						
Number of obs		836	164	BW type	=	mserd
Eff. Number of obs		129	83	Kernel	=	Triangular
Order est. (p)		1	1	VCE method	=	NN

```

Order bias (q) |          2          2
BW est. (h) |    0.413    0.413
BW bias (b) |    0.658    0.658
rho (h/b) |    0.627    0.627

```

Outcome: y. Running variable: xc.

Method	Coef.	Std. Err.	z	P> z	[95% Conf. Interval]	
Conventional	.41482	.47553	0.8723	0.383	-.517206	1.34685
Robust	-	-	0.9336	0.350	-.581835	1.64038

```

.      local bandwidth = e(h_1)

.      display 'bandwidth'
.41251756

.      // pass through bandwidth to get local RD plot
.
.      rdplot y xc if abs(xc) <= 'bandwidth', c(0) p(1) h('bandwidth') ///
>          kernel(triangular) graph_options(legend(position(6)) name(ballou, replace))


```

RD Plot with evenly spaced mimicking variance number of bins using spacings estimato rs.

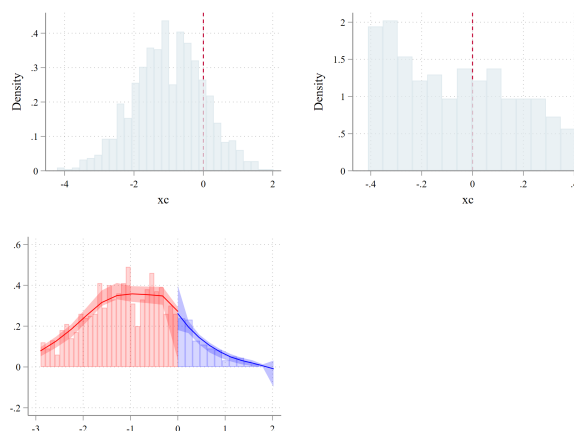
Cutoff c = 0   Left of c Right of c			Number of obs =	212
			Kernel	= Triangular
Number of obs	129	83		
Eff. Number of obs	129	83		
Order poly. fit (p)	1	1		
BW poly. fit (h)	0.413	0.413		
Number of bins scale	1.000	1.000		

Outcome: y. Running variable: xc.

	Left of c	Right of c
Bins selected	10	8
Average bin length	0.041	0.051
Median bin length	0.041	0.051
IMSE-optimal bins	7	5
Mimicking Var. bins	10	8
Rel. to IMSE-optimal:		
Implied scale	1.429	1.600
WIMSE var. weight	0.255	0.196
WIMSE bias weight	0.745	0.804



- (e) Check for manipulation in the running variable  $xc$  in two ways: by inspection using histogram, and using `rddensity`. Do you expect manipulation here? What does the test conclude?



```
.      rddensity xc, c(0) plot graph_opt(name(denstest,replace) legend(off))
Computing data-driven bandwidth selectors.
```

RD Manipulation test using local polynomial density estimation.

c =	0.000	Left of c	Right of c	Number of obs =	1000
-----+-----					
Number of obs	836	164	Model	=	unrestricted
Eff. Number of obs	313	134	BW method	=	comb
Order est. (p)	2	2	Kernel	=	triangular
Order bias (q)	3	3	VCE method	=	jackknife
BW est. (h)	0.966	0.858			

Running variable: xc.

-----		
Method	T	P> T

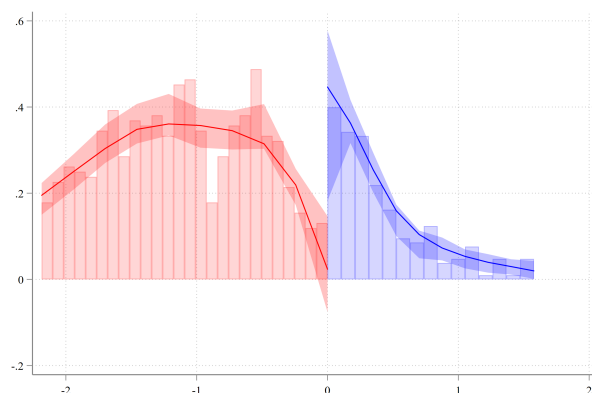
```
-----+-----
Robust |      1.4211      0.1553
-----+-----
```

The formal density test does not detect any manipulation around the cut-point. The test statistic is 1.4211 with a p-value of 0.1553. We cannot reject the null hypothesis of no manipulation. We would not expect any in this case, since we generated the data with no manipulation present.

- (f) Now modify the data a bit to introduce manipulation in  $x$ . Try the syntax below and explain in words what the first line is doing. (Create a new centered  $x$  variable  $xcm$ ) Then, re-do the `rddensity` test. Does it detect the manipulation?

```
gen xcm = x
replace xcm = xcm + .4 if xcm < 1 & xcm > .65 & e > 0
gen xmc = xcm-1
```

The code above is manipulating values of  $x$  between 0.65 and 1, giving them an additional 0.4 to put them over the threshold. In this case the density test is clear in showing the manipulation. The test statistic is 2.9 with a p-value of 0.0036, so we can reject the null hypothesis of no manipulation.



```
.      rddensity xmc, c(0) plot graph_opt(name(denstest2,replace) legend(off))
Computing data-driven bandwidth selectors.
```

RD Manipulation test using local polynomial density estimation.

c =	0.000	Left of c	Right of c	Number of obs =	1000
-----+-----					
Number of obs	784	216	Model	=	unrestricted
Eff. Number of obs	201	153	BW method	=	comb
Order est. (p)	2	2	Kernel	=	triangular
			VCE method	=	jackknife

```

Order bias (q) |          3          3
BW est. (h) |    0.728    0.526

```

Running variable: xmc.

```

-----
Method |      T      P>|T|
-----+-----
Robust |  2.9102  0.0036
-----

```

- (g) Now that we know there is manipulation, try estimating the nonparametric RD using `rdrobust` (use same specs as part d). What does this yield?

For this step it is worth thinking about how  $y$  should be changed, if at all. If we assume that cases manipulated into the treatment group get the same effect from being exposed to the treatment, then we should add the 0.5 to these cases (as below). One could also leave the original  $y$ 's intact, but this would be assuming no treatment effect for these manipulated cases. The RD estimate of the treatment effect is too small in this case: -0.23. This is likely because cases that would have (without manipulation) not been treated were in fact treated. These cases had somewhat lower potential outcomes than other non-manipulated cases just above the cutoff.

```

.      gen y2=y

.      replace y2=y+0.5 if (x<1 & x>0.65 & e>0)
(52 real changes made)

.      rdrobust y2 xmc, c(0) p(1) bwselect(mserd) kernel(triangular)

```

Sharp RD estimates using local polynomial regression.

```

Cutoff c = 0 | Left of c  Right of c          Number of obs =      1000
-----+-----
Number of obs |      784      216          BW type      =      mserd
Eff. Number of obs |      44      111          Kernel        = Triangular
Order est. (p) |       1       1          VCE method     =      NN
Order bias (q) |       2       2
BW est. (h) |    0.308    0.308
BW bias (b) |    0.512    0.512
rho (h/b) |    0.602    0.602

```

Outcome: y2. Running variable: xmc.

```

-----
Method |   Coef.   Std. Err.   z   P>|z|   [95% Conf. Interval]
-----+-----
Conventional | -.22507   .53709   -0.4191  0.675   -1.27774   .8276
Robust |      -      -   -0.0206  0.984   -1.2555   1.22936
-----

```



**Example 3.** This example, based on an example created by Celeste Carruthers, also uses simulated data to estimate the effect of participation in a gifted and talented (G&T) program.

- (a) Generate 10,000 student observations. The data will include a measure of students' "true ability,"  $trueability \sim N(50, 4)$ , and their 3rd grade test score, which is a noisy measure of their true ability  $grade3test = trueability + u$  where  $u \sim N(0, 1)$ . To add a bit of realism, we will round test scores to the nearest 0.25 to create a discrete scale.

```
clear
set seed 195423
set obs 10000
gen id=_n
gen trueability = 50 + 4*rnormal()
gen grade3test = trueability + rnormal()
replace grade3test = round(grade3test, 0.25)
```

- (b) Suppose 3rd graders scoring at or above 56 are eligible for the G&T program. Create a treatment assignment variable re-centered at zero, and a "gap" variable that contains the distance between the running variable and the cut score.

```
gen above56 = (grade3test>=56)
gen gap = grade3test-56
```

- (c) Assume perfect compliance. Create an indicator variable for G&T participation *inGT*, that equals one for treated students and zero otherwise. What proportion of students are treated? Try estimating an OLS regression for G&T participation where *inGT* is regressed on the *gap* and the threshold indicator *above56*. What happens and why?

**7.59% of students participated in G&T. When you regress the treatment *inGT* on the *gap* and threshold indicator *above56*, Stata cannot produce estimates. This is because *above56* perfectly determines the outcome *inGT*. (It is a strict, not fuzzy, discontinuity).**

```
. gen inGT=(above56==1)
```

```
. sum inGT
```

Variable	Obs	Mean	Std. Dev.	Min	Max
inGT	10,000	.0759	.2648513	0	1

```
. reg inGT gap above56
```

Source	SS	df	MS	Number of obs	=	10,000
Model	701.3919	2	350.69595	F(2, 9997)	=	.
				Prob > F	=	.



```
.      reg grade4test c.gap##i.above56
```

Source	SS	df	MS	Number of obs	=	10,000
Model	187161.733	3	62387.2445	F(3, 9996)	=	32218.65
Residual	19355.9606	9,996	1.9363706	Prob > F	=	0.0000
				R-squared	=	0.9063
				Adj R-squared	=	0.9062
Total	206517.694	9,999	20.6538348	Root MSE	=	1.3915

grade4test	Coef.	Std. Err.	t	P> t	[95% Conf. Interval]	
gap	.939221	.0040802	230.19	0.000	.9312229	.9472191
1.above56	3.02236	.0798731	37.84	0.000	2.865792	3.178927
above56#c.gap						
1	.0046223	.0313781	0.15	0.883	-.0568851	.0661296
_cons	60.61455	.0306168	1979.78	0.000	60.55453	60.67456

```
.      rdrobust grade4test gap, c(0) p(1) kernel(triangular)
```

Mass points detected in the running variable.

Sharp RD estimates using local polynomial regression.

Cutoff c = 0	Left of c	Right of c	Number of obs	=	10000
			BW type	=	mserd
Number of obs	9241	759	Kernel	=	Triangular
Eff. Number of obs	928	525	VCE method	=	NN
Order est. (p)	1	1			
Order bias (q)	2	2			
BW est. (h)	2.190	2.190			
BW bias (b)	3.633	3.633			
rho (h/b)	0.603	0.603			
Unique obs	80	33			

Outcome: grade4test. Running variable: gap.

Method	Coef.	Std. Err.	z	P> z	[95% Conf. Interval]	
Conventional	3.2231	.15821	20.3725	0.000	2.91305	3.53322
Robust	-	-	17.4736	0.000	2.90155	3.63471


Estimates adjusted for mass points in the running variable.

- (f) Drop the existing *inGT* and *grade4test* variables and re-create them assuming a “fuzzy” GT treatment that increases smoothly with grade 3 test scores and then jumps discontinuously (by about 70 percentage points) at the cut score. This might arise if G&T placement is dependent on the grade 3 test score as well as other factors (e.g., parental

input, teacher recommendation). Use the syntax below. Now what proportion of students are treated, overall? Below the cutoff? Above? Use `binscatter` to visualize the relationship between treatment and the grade 3 score.

```
drop inGT grade4test
gen inGT=round(-.77+.007*grade3test+0.7*above56+runiform())
gen grade4test = round(trueability + 5 + rnormal() + (3*inGT), 0.25)
```

13.3% of students are treated, overall. Below the cutoff, 7.6% of students are selected for the G&T program. Above the cutoff, 83.4% are selected. The discontinuous jump in treatment status is seen in the binscatter below.



- (g) As in (c), estimate a regression for G&T placement where *inGT* is regressed on the *gap* and the threshold indicator *above56*. Interpret your results. (Try estimating this in two ways: first assuming the slope is constant on either side of the cutoff, and then allowing the slope to change). For later use, use the `predict` command to get predicted values for treatment (placement in G&T) given the 3rd grade score. Call this variable *predGT*.

**Results below.** The first regression tells us that the probability of selection for G&T increases with *gap* (the student's score minus 56). There is also a discontinuous jump in the probability of selection at the cut score, of 70.3 percentage points.

```
.      reg inGT gap i.above56
```

Source	SS	df	MS	Number of obs	=	10,000
Model	408.373259	2	204.186629	F(2, 9997)	=	2730.15
Residual	747.671141	9,997	.074789551	Prob > F	=	0.0000
Total	1156.0444	9,999	.115616002	R-squared	=	0.3533
				Adj R-squared	=	0.3531
				Root MSE	=	.27348

inGT	Coef.	Std. Err.	t	P> t	[95% Conf. Interval]	
gap	.0066524	.0007951	8.37	0.000	.0050939	.0082109
1.above56	.7026519	.012272	57.26	0.000	.6785962	.7267075
_cons	.1198432	.0059775	20.05	0.000	.1081262	.1315602

```
. reg inGT c.gap##i.above56
```

Source	SS	df	MS	Number of obs	=	10,000
Model	408.418692	3	136.139564	F(3, 9996)	=	1820.23
Residual	747.625708	9,996	.074792488	Prob > F	=	0.0000
Total	1156.0444	9,999	.115616002	R-squared	=	0.3533
				Adj R-squared	=	0.3531
				Root MSE	=	.27348

inGT	Coef.	Std. Err.	t	P> t	[95% Conf. Interval]	
gap	.0065711	.0008019	8.19	0.000	.0049992	.008143
1.above56	.695023	.0156977	44.28	0.000	.6642524	.7257936
above56#c.gap						
1	.0048064	.0061668	0.78	0.436	-.0072818	.0168946
_cons	.1193058	.0060172	19.83	0.000	.1075109	.1311008

```
. // get predicted treatment for later use
. predict predGT
(option xb assumed; fitted values)
```

- (h) Re-estimate the RD models from part (e) assuming a linear relationship with the running variable. Assume the discontinuity is “sharp,” even though we know otherwise. How does the estimated treatment effect differ from the known treatment effect of 3 points?

**Results below. The estimated treatment effect is smaller, at 2.0 versus the known 3 points. This is not surprising: when the discontinuity is fuzzy, the difference in outcomes around the cutoff will be smaller, since not all above the cut score were treated, and some of those below the cut score were treated.**

```
. reg grade4test gap i.above56
```

Source	SS	df	MS	Number of obs	=	10,000
Model	180812.138	2	90406.0691	F(2, 9997)	=	35035.27
Residual	25796.5606	9,997	2.58043019	Prob > F	=	0.0000
				R-squared	=	0.8751

```
-----+-----
Total | 206608.699    9,999  20.6629362  Adj R-squared = 0.8751
Root MSE = 1.6064
```

```
-----+-----
grade4test |      Coef.   Std. Err.      t    P>|t|     [95% Conf. Interval]
-----+-----
      gap |   .9617267   .0046702   205.93  0.000   .9525722   .9708812
1.above56 |   2.06733   .0720845    28.68  0.000   1.92603   2.208631
   _cons |   61.0054   .0351109   1737.50  0.000   60.93658   61.07423
-----+-----
```

```
.      reg grade4test c.gap##i.above56
```

```
-----+-----
Source |      SS          df           MS       Number of obs = 10,000
-----+-----
Model | 180814.673            3   60271.5578       F(3, 9996) = 23357.13
Residual | 25794.0253         9,996    2.5804347       Prob > F = 0.0000
-----+-----
Total | 206608.699         9,999   20.6629362       R-squared = 0.8752
Adj R-squared = 0.8751
Root MSE = 1.6064
```

```
-----+-----
grade4test |      Coef.   Std. Err.      t    P>|t|     [95% Conf. Interval]
-----+-----
      gap |   .9611196   .0047102   204.05  0.000   .9518867   .9703525
1.above56 |   2.010342   .0922047    21.80  0.000   1.829602   2.191082
above56#c.gap |
1 |   .0359043   .0362225     0.99  0.322   -.0350991   .1069077
   _cons |   61.00139   .0353438   1725.95  0.000   60.93211   61.07067
-----+-----
```

```
.      rdrobust grade4test gap, c(0) p(1) kernel(triangular)
Mass points detected in the running variable.
```

Sharp RD estimates using local polynomial regression.

```
-----+-----
Cutoff c = 0 | Left of c  Right of c           Number of obs = 10000
-----+-----
Number of obs |      9241      759           BW type = mserd
Eff. Number of obs |      781      487           Kernel = Triangular
Order est. (p) |        1        1           VCE method = NN
Order bias (q) |        2        2
BW est. (h) |    1.995    1.995
BW bias (b) |    3.275    3.275
rho (h/b) |    0.609    0.609
Unique obs |        80       33
```

Outcome: grade4test. Running variable: gap.

```
-----+-----
Method |      Coef.   Std. Err.      z    P>|z|     [95% Conf. Interval]
```



Conventional		.74416	.03788	19.6451	0.000	.66992	.818409
Robust		-	-	17.9063	0.000	.674327	.84009

Treatment effect estimates. Outcome: grade4test. Running variable: gap. Treatment St  
> atus: inGT.

Method		Coef.	Std. Err.	z	P> z	[95% Conf. Interval]
Conventional		3.0414	.21207	14.3415	0.000	2.62573 3.45702
Robust		-	-	12.7669	0.000	2.58462 3.52212

Estimates adjusted for mass points in the running variable.

```
.      // then 2SLS (second stage shown here using predicted GT)
.      reg grade4test gap predGT
```

Source		SS	df	MS	Number of obs	=	10,000
Model		180813.742	2	90406.8711	F(2, 9997)	=	35037.76
Residual		25794.9566	9,997	2.58026974	Prob > F	=	0.0000
Total		206608.699	9,999	20.6629362	R-squared	=	0.8752
					Adj R-squared	=	0.8751
					Root MSE	=	1.6063

grade4test		Coef.	Std. Err.	t	P> t	[95% Conf. Interval]
gap		.942128	.0050712	185.78	0.000	.9321874 .9520686
predGT		2.943022	.1025765	28.69	0.000	2.741951 3.144093
_cons		60.65253	.0434644	1395.45	0.000	60.56734 60.73773

```
.      // ivregress is tricky here since "above56" is an exogenous instrument
.      // conditional on grade3test. Can't use above56 both as the instrument
.      // and in the interaction with "gap". So manually created an interaction
.      // below
.
.      gen gapabove = gap*above56
.
.      gen gapbelow = gap*(1-above56)
.
.      // note gapbelow and gapabove estimate slope on gap below and above
.      ivregress 2sls grade4test (inGT=above56) gapbelow gapabove , first
```

First-stage regressions

Number of obs	=	10,000
F( 3, 9996)	=	1820.23
Prob > F	=	0.0000
R-squared	=	0.3533



Adj R-squared = 0.3531  
Root MSE = 0.2735

inGT	Coef.	Std. Err.	t	P> t	[95% Conf. Interval]	
gapbelow	.0065711	.0008019	8.19	0.000	.0049992	.008143
gapabove	.0113775	.0061145	1.86	0.063	-.0006081	.0233631
above56	.695023	.0156977	44.28	0.000	.6642524	.7257936
_cons	.1193058	.0060172	19.83	0.000	.1075109	.1311008

Instrumental variables (2SLS) regression      Number of obs = 10,000  
Wald chi2(3) = 93654.38  
Prob > chi2 = 0.0000  
R-squared = 0.9066  
Root MSE = 1.3895

grade4test	Coef.	Std. Err.	z	P> z	[95% Conf. Interval]	
inGT	2.892483	.1147519	25.21	0.000	2.667573	3.117392
gapbelow	.9421128	.0043867	214.77	0.000	.933515	.9507106
gapabove	.9641147	.0319592	30.17	0.000	.9014757	1.026754
_cons	60.6563	.0379859	1596.81	0.000	60.58185	60.73075

Instrumented: inGT

Instruments: gapbelow gapabove above56

- (j) Finally, do a manipulation test using `rddensity`. What does it find? Would you expect to see evidence of manipulation here?

**Results are below. We cannot reject the null hypothesis of no manipulation. This is the correct decision, since we know from generating the data that 3rd grade test scores were not manipulated.**

```
. rddensity gap, c(0) plot graph_opt(name(fuzzy, replace) legend(off))
Computing data-driven bandwidth selectors.
```

Point estimates and standard errors have been adjusted for repeated observations.  
(Use option `nomasspoints` to suppress this adjustment.)

RD Manipulation test using local polynomial density estimation.

c =	0.000	Left of c	Right of c	Number of obs =	10000
-----+-----				Model =	unrestricted
Number of obs	9241	759		BW method =	comb
Eff. Number of obs	4153	737		Kernel =	triangular
Order est. (p)	2	2		VCE method =	jackknife
Order bias (q)	3	3			

BW est. (h) |        5.750        5.750

Running variable: gap.

Method	T	P> T
Robust	1.5199	0.1285

